

Retail Sales Analytics – SQL Data Layer & Governance Framework

Project Objective

This project extends the Retail Sales Power BI dashboard by implementing a structured SQL data layer with controlled transformations, referential integrity, and certified revenue logic before exposing data to reporting tools.

The goal was to simulate enterprise-level architecture:

Raw Data → Clean Star Schema → Certified Metric Layer → Power BI Reporting

Architecture Overview

1. Staging Layer (stg_)

- Direct import of raw CSV data
- No transformations applied
- Used for audit and validation checks

Tables:

- stg_Customers
- stg_Products
- stg_Stores
- stg_Transactions

2. Clean Layer (Star Schema)

Implemented structured dimensional modeling:

Dimension Tables

- dim_Customers
- dim_Products
- dim_Stores

Fact Table

- fact_Transactions

Key Design Decisions

- Business keys stored as NVARCHAR(50) (e.g., C001, P002)
- Primary Keys enforced on all dimensions
- Foreign Keys enforced in fact table
- CHECK constraints applied on:
 - Quantity (> 0)
 - Discount (0–1 range)
 - UnitPrice and CostPrice (≥ 0)

This ensures referential integrity and data quality control at the database level.

Data Quality Validation

Before loading into clean tables, the following checks were performed:

- NULL checks on primary identifiers
- Duplicate key validation
- Negative quantity detection
- Discount range validation
- Unit price vs cost price validation

Example validation query:

```
SELECT CustomerID, COUNT(*)
FROM stg_Customers
GROUP BY CustomerID
HAVING COUNT(*) > 1;
```

All staging data passed integrity checks prior to transformation.

Certified Revenue Definition

To prevent metric drift, revenue logic was centralized in a SQL view:

vw_Certified_Revenue

Revenue Formula

$$\text{Net Revenue} = \text{Quantity} \times \text{UnitPrice} \times (1 - \text{Discount})$$

Gross Profit Formula

$$\text{Gross Profit} = \text{Net Revenue} - (\text{Quantity} \times \text{CostPrice})$$

SQL Implementation:

```
CREATE OR ALTER VIEW vw_Certified_Revenue AS
SELECT
    t.TransactionID,
    t.Date,
    t.CustomerID,
    t.ProductID,
    t.StoreID,
    t.Quantity,
    p.UnitPrice,
    p.CostPrice,
    t.Discount,
    (t.Quantity * p.UnitPrice * (1 - t.Discount)) AS CertifiedRevenue,
    ((t.Quantity * p.UnitPrice * (1 - t.Discount))
     - (t.Quantity * p.CostPrice)) AS GrossProfit
FROM fact_Transactions t
JOIN dim_Products p
    ON t.ProductID = p.ProductID;
```

This ensures:

- Single source of truth for revenue
- Consistency across reports
- Prevention of alternate revenue calculations

Analytical SQL Demonstrations

To validate business logic prior to visualization, the following analytical queries were implemented:

Year-wise Revenue

```
SELECT  
    YEAR(Date) AS Year,  
    SUM(CertifiedRevenue) AS TotalRevenue,  
    SUM(GrossProfit) AS TotalProfit  
FROM vw_Certified_Revenue  
GROUP BY YEAR(Date)  
ORDER BY Year;
```

Top Products by Revenue (Ranking)

```
SELECT TOP 5  
    p.ProductName,  
    SUM(v.CertifiedRevenue) AS TotalRevenue,  
    RANK() OVER (ORDER BY SUM(v.CertifiedRevenue) DESC) AS RevenueRank  
FROM vw_Certified_Revenue v  
JOIN dim_Products p ON v.ProductID = p.ProductID  
GROUP BY p.ProductName  
ORDER BY TotalRevenue DESC;
```

Month-over-Month Revenue (CTE + Window Function)

```
WITH MonthlyRevenue AS (  
    SELECT  
        YEAR(Date) AS YearNum,  
        MONTH(Date) AS MonthNum,  
        SUM(CertifiedRevenue) AS TotalRevenue  
    FROM vw_Certified_Revenue  
    GROUP BY YEAR(Date), MONTH(Date)  
)  
SELECT  
    YearNum,  
    MonthNum,  
    TotalRevenue,  
    LAG(TotalRevenue) OVER (ORDER BY YearNum, MonthNum) AS PreviousMonthRevenue,  
    TotalRevenue  
        - LAG(TotalRevenue) OVER (ORDER BY YearNum, MonthNum) AS RevenueChange  
FROM MonthlyRevenue  
ORDER BY YearNum, MonthNum;
```

Validation Against Power BI

All SQL aggregates were validated against Power BI DAX measures to ensure full consistency between:

- SQL data layer
- Semantic layer
- Dashboard outputs

No discrepancies were found.

Key Outcomes

- Implemented layered SQL architecture (Raw → Clean → Certified)
- Enforced referential integrity via PK/FK constraints
- Centralized revenue governance logic
- Demonstrated CTEs and window functions
- Validated analytical outputs prior to visualization

This project demonstrates full ownership of the data lifecycle — from ingestion and validation to governed metric design and analytical reporting.